


- 1)  $s_1(v) r_1(x) x_1(y) w_3(x) x_1(y) w_3(y) x_1(z) s_1(v) u_3(x) s_1(x) r_1(x) u_3(y) s_1(y) r_2(y) u_1(x) u_1(v) x_1(x) w_2(x)$   
 $w_3(z) u_3(z) s_1(z) r_2(z) s_1(v) r_2(v) r_3(v) u_3(v) x_1(z) u_3(v) x_1(v) w_4(v)$  ! Now  $T_4$  and  $T_2$  both want access to  $z \rightarrow$  Not 2PL
- $T_4 < T_2$  for the last write of  $z$   
 $T_2 < T_4$  because must read from "default"  $v \rightarrow$  impossible  $\rightarrow$  Not view-serializable
  - Not ACID cause  $r_1(x)$  reads from  $T_3$  not already committed  
 $\Rightarrow$  Not strict  $\Rightarrow$  Not rigorous
  - Not recoverable cause  $T_1$  commits before  $T_3$

2) By definition if a schedule is accepted by a timestamped scheduler then it is conflict-serializable  $\Rightarrow$  view-serializable.

Thomas Rules just ignore outdated writes and the obtained schedule is conflict serializable

- 3) Sorted: entries are sorted  
 Primary: search key includes primary key of relation  
 Dense: one data entry for each data record of data file  
 Clustered: sorted according some criteria of data file



To find the given  $k$  we can use binary search.

- 4) Ordered file on code  
 $1K/20 = 50$  fields per page  $\rightarrow 5$  tuples per page  $\rightarrow 100K/5 = 200,000$  pages  
 • Since we have to show all pages we do 200,000 p.a.  
 • To access the first needed page (code  $\geq 30$ ) we use binary search:  $\log_2 200,000 + \# \text{page needed to store page with code} \geq 30$  and  $\leq 50$
- 5)  $R$  without duplicates must fit  $M-1$  buffers: read one page at time and keep track of already seen records in hash/tree in  $M-1$  buffers  
 • Read  $R$  and create  $M-1$  sorted sublists of  $M$  pages. Load one page for each of them check duplicates  
 $B(R)/M \leq M-1 \sim B(R)/M + 1 \leq M-1 = B(R) \leq M \cdot (M-2) \sim B(R) \leq (M-2)^2 \Rightarrow$  we need  $\sqrt{B(R)} + 2$  buffers frames